

EC 535: Introduction to Embedded Systems

Eshed Ohn-Bar
1/22/2025

TA Office Hours:

4-5PM, Tu/Th, PHO 307



Homework ▼



In-Class Exercises ▼

Challenges in embedded system design

- How much hardware do we need?
 - How many processors? How big are they? How much memory?
- How do we meet performance requirements?
 - What's in hardware? What's in software?
 - Faster hardware or cleverer software?

Optimization by doing things in a certain order

constraint to make system responsive in real-time, need to do optimization on SW and/or HW
- How do we minimize cost and power? → mechanisms to preserve energy
 - Turn off unnecessary logic? Reduce memory accesses?
- How do we ship in time?
 - Buy solution? From scratch? Off-the-shelf chips? IP-reuse?

↳ different ways to speed up process is outsource parts

Automotive industry

AI solutions bottlenecked into CUDA and NVIDIA

optimizing codes req. optimize CUDA code → very hard...?

Challenges (cont'd)

testing ABS is hard to do
- Simulating hardware

• How do we know that it fully works? *extensive testing*

- Is the specification correct?
- Does the implementation meet the spec?
- How do we test for real-time characteristics?
- How do we test on real data?

} specifications to verify
Functional aspects to
verify and see if we meet
them in testing

• How do we locate the problem if it doesn't work?

- Observability, controllability?

- hard to debug embedded systems
because no monitoring, various debugging tools but more difficult
than general purpose systems

difficult to do real-world testing for embedded systems

★ HW in-the-loop simulation

↳ phys controller testing in loop

- simulated inputs to system in closed environment

Embedded System Designers

- Expertise with both **software and hardware** is needed to optimize design metrics
 - Not just a hardware or software expert.
 - A designer must be comfortable with various technologies in ^{domain knowledge} order to choose the best for a given application and constraints.
 - A designer must be able to **communicate** with teammates of various backgrounds.

Need to approach in all aspects

Embedded Systems – Design methodologies

- A procedure for designing a system.
- Understanding your methodology helps you ensure you didn't skip anything.
- Compilers, software engineering tools, computer-aided design (CAD) tools, etc., can be used to:
 - help automate methodology steps;
 - keep track of the methodology itself.

automate parts of the design process
where it can check system automatically.
★ Track design process

↳ Git!! version control

Top-down vs. bottom-up

- Top-down design:

- start from most abstract description;
- work to most detailed.

What the system needs to do

"I want a car to get me to the destination"


↳ think about what customer wants, often want to do things but we don't have avail. resources

- Bottom-up design:

- work from small components to big system.

→ figure out what we have, integrate and build up to what we can achieve

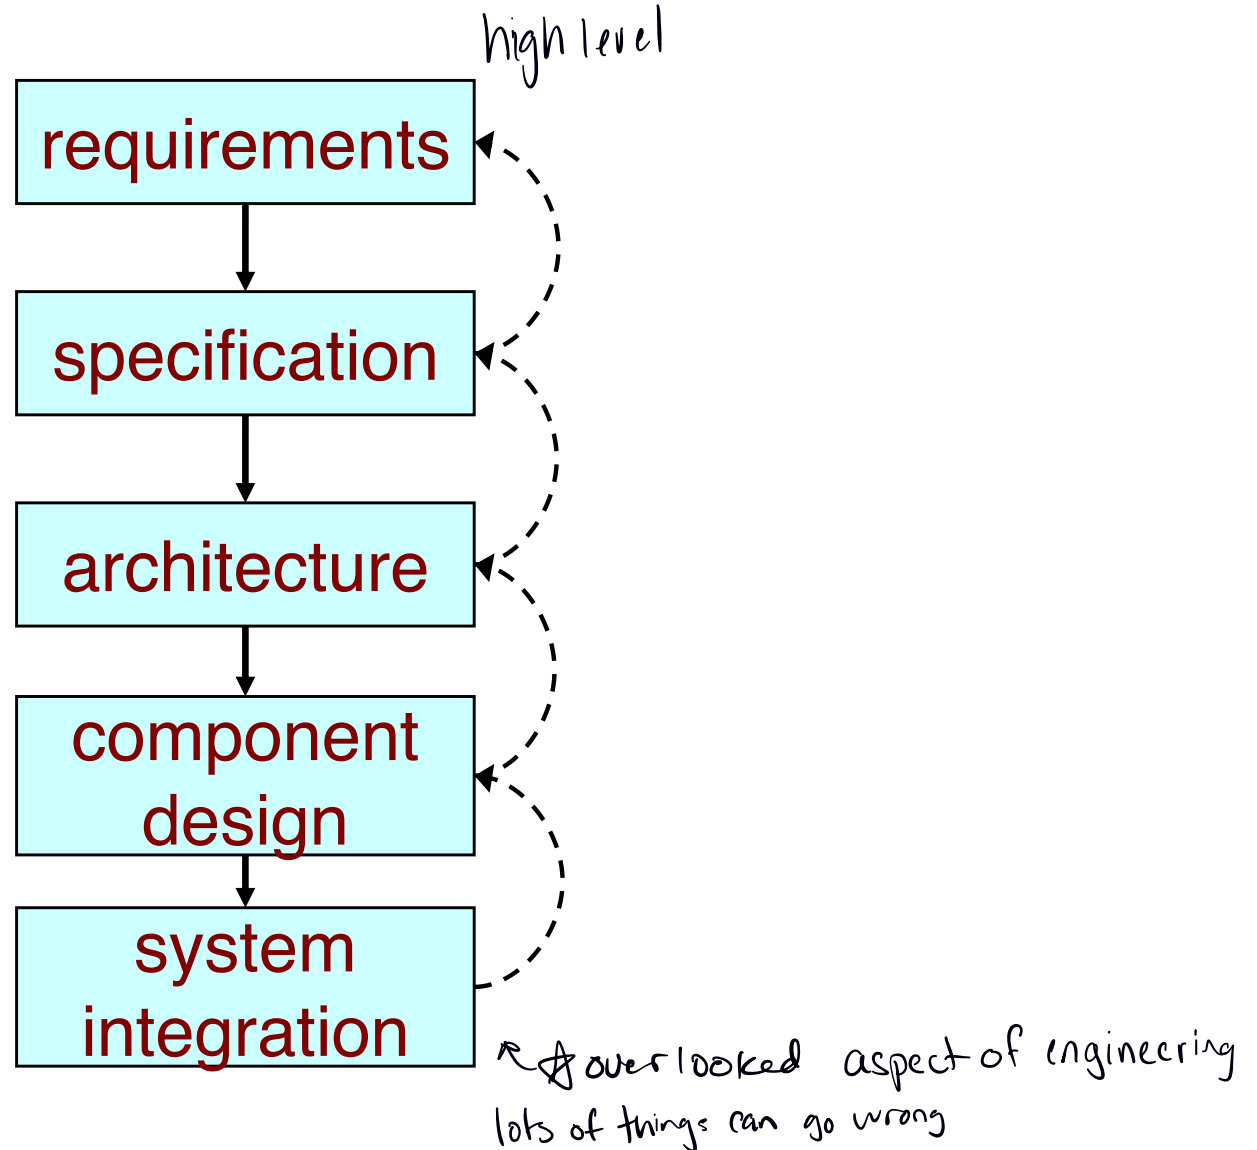
- Real design uses both techniques, more or less.



★ combine approaches

★ Apple display: design from scratch

Levels of abstraction



Requirements

*expectations in language
for system*

- Plain language description of what the user wants and expects to get.
- Functional requirements:
 - output as a function of input.
- Non-functional requirements:
 - time required to compute output;
 - size, weight, etc.;
 - power consumption;
 - reliability;
 - etc.
- May be developed in several ways:
 - talking directly to customers;
 - talking to marketing representatives;
 - providing prototypes to users for comment.

**What types of requirements require less
redesigns*

Specification

- A more precise description of the system:
 - should not imply a particular architecture;
 - provides input to the architecture design process.
- May include functional and non-functional elements.
 - system behavior
input → output
 - performance constraints
- May be executable or may be in mathematical form for proofs.
 - ↓
* formal methods?

Architecture design

- What major components do we need to satisfy the specification?
- Hardware components:
 - CPUs, peripherals, etc.
- Software components:
 - Major programs and their operations.
- Must take into account functional and non-functional specifications.

Designing hardware and software components

algorithm to adjust temp
software-based

- Must spend time architecting the system before you start coding.
- Some components are ready-made, some can be modified from existing designs, others must be designed from scratch.

SW - flexible part of system

HW - inflexible but great rewards for optimized

GPU: markets have dominated by Nvidia,
don't need to modify too much

Some bought unchanged,
power management algo is new,
modified code

} → hard task to do new products

System integration


★ Undervalued
hard to do

- Put together the components.
 - Many bugs appear only at this stage.
- Have a plan for integrating components to uncover bugs quickly, test as much functionality as early as possible.

Assumptions on HW and SW

- **Hardware:** In this course, *hardware* means *single-clock synchronous digital circuits* that are created using word-level combinational and sequential building blocks. These circuits can be modeled with building blocks such as *registers, adders, multiplexers*, ... Perform various ops
↳ basic building blocks of embedded system

- There are well-known circuit symbols for these building blocks.
- Using these circuit elements, schematics can be created.
- Schematics quickly become incomprehensible with increasing circuit complexity, hardware designers very often use a textual representation in practice.



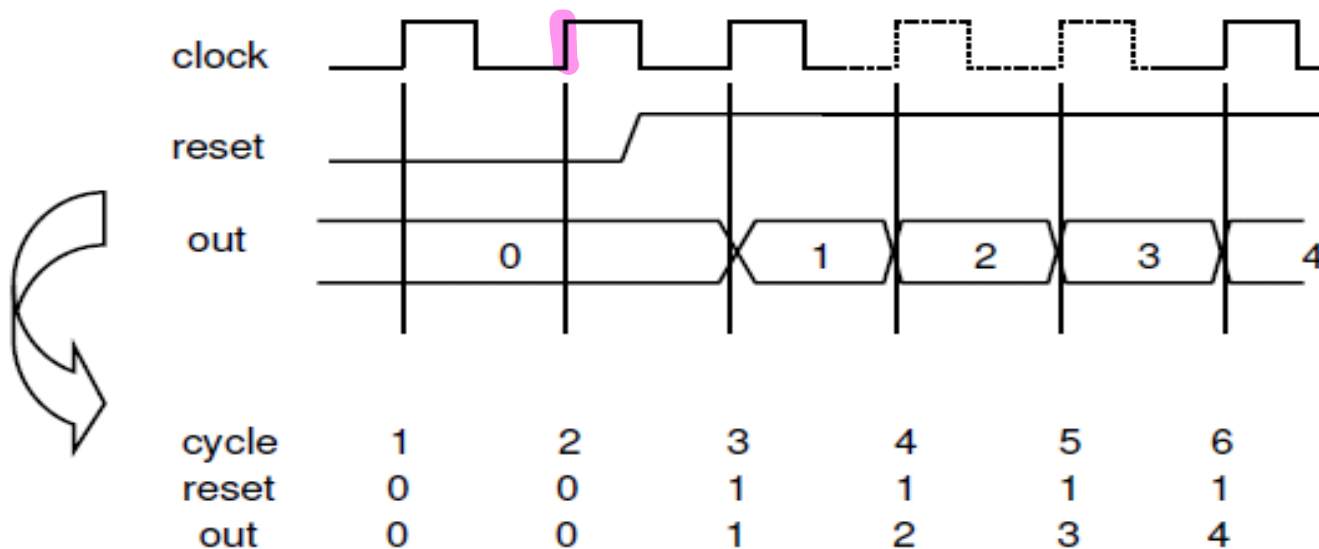

Other Types of HW?

★ course is synchronous single-clock
easier to work with

- Asynchronous hardware
- Dynamic logic
- Multi-phase clocked hardware
- etc.

→ We'll focus on synchronous single-clock hardware

Synchronous HW Model

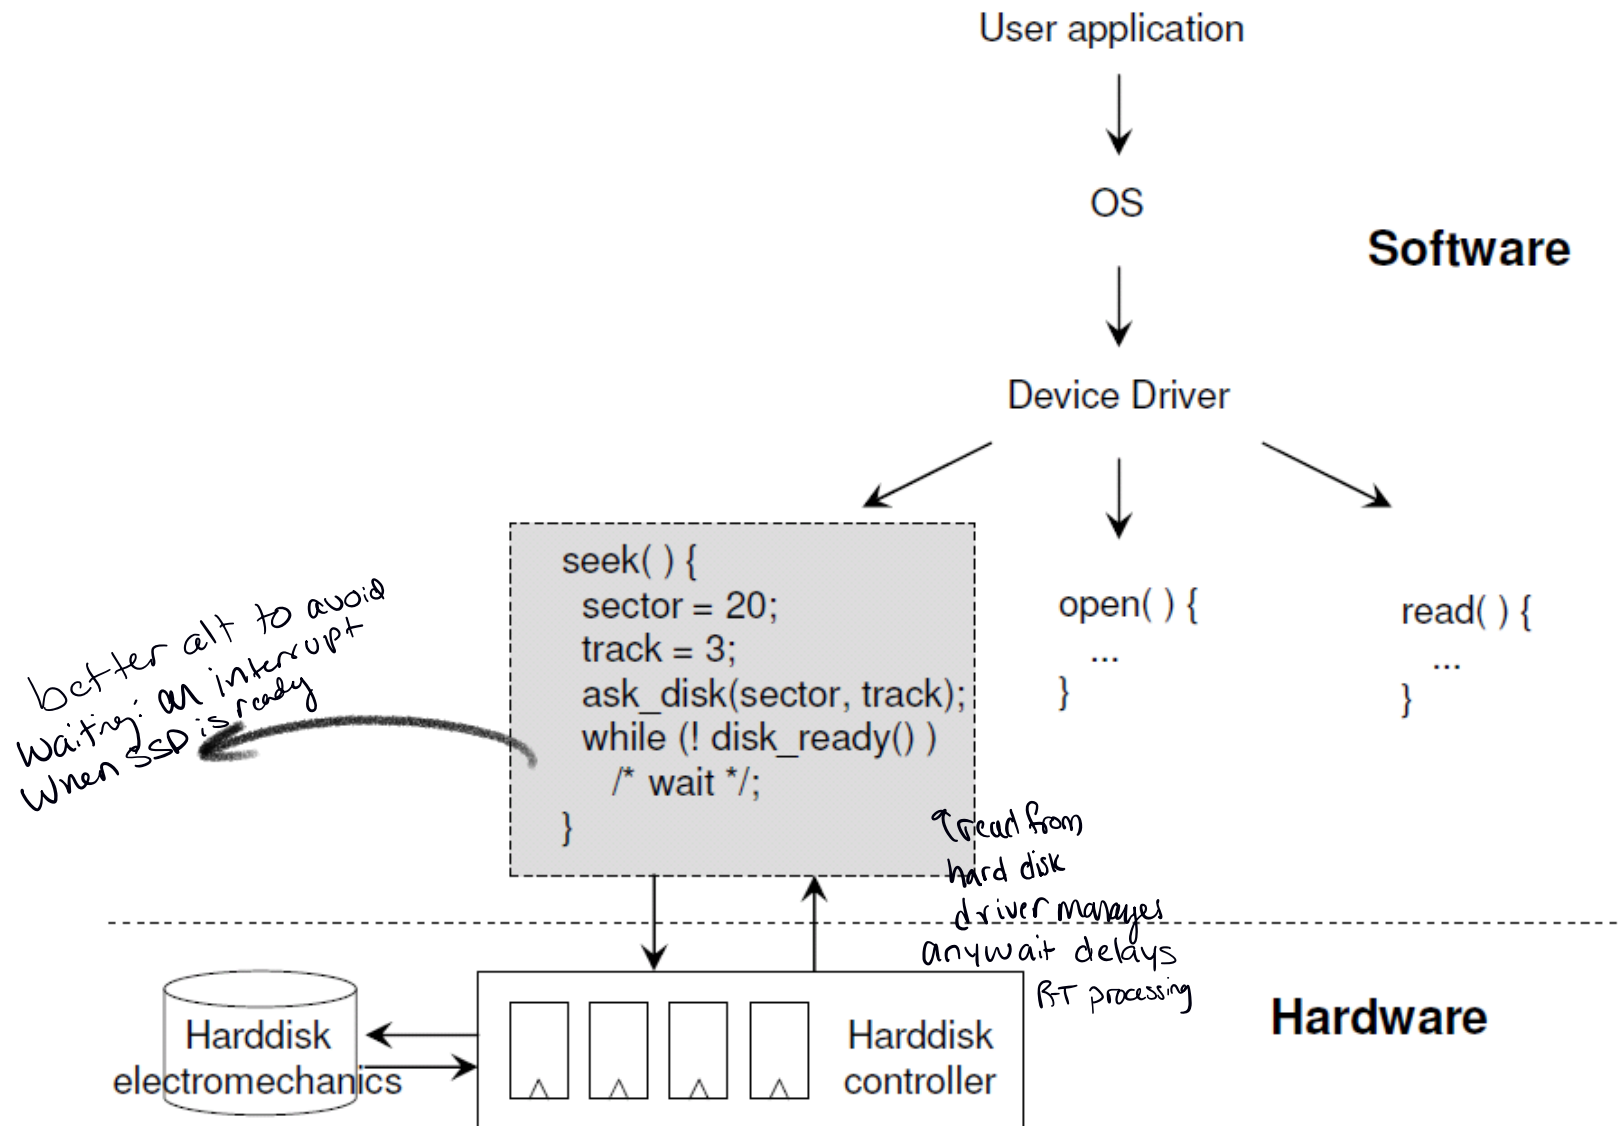


Software

- With software we [mostly] mean single-thread sequential programs, such as those represented by a C program or an assembly program.
*One processor one task
One op. at any time : issue for real-time components*
- We make abstractions of underlying mechanisms to enhance concurrency (such as an operating system).
*allows multiple things
to 'run' at same time*

Device Driver for Hard Disk

interface between SW+HW



Comparing Hardware and Software

HW w/ time and clock

- ^{HW}Parallel versus ^{SW}sequential operation
- Spatial versus temporal (time-wise) decomposition
- Flexibility
- Control processing versus data processing
- Modeling and implementation
- Intellectual property reuse

Main takeaway


Concurrency and Parallelism

OS can manage multiple tasks concurrently

- Concurrency is the ability to execute simultaneous operations (because these operations are completely independent).
SW sequential but can incorporate concurrency
 - Relates to an application model.
- Parallelism is the ability to execute simultaneous operations because the operations can run on different processors or circuit elements.
HW usually parallel
 - Relates to the implementation of the application model.
- HW: Parallel?
- SW: Sequential, concurrent, parallel?

An example parallel processor: Connection Machine (CM)

*network of nodes, running
own C, can be parallelized*



*Danny Hillis,
1980s, MIT*

- How to write programs for such a parallel architecture?

- Individual C programs for each node?
 - Complexity
 - Scalability
- Concurrent algorithms

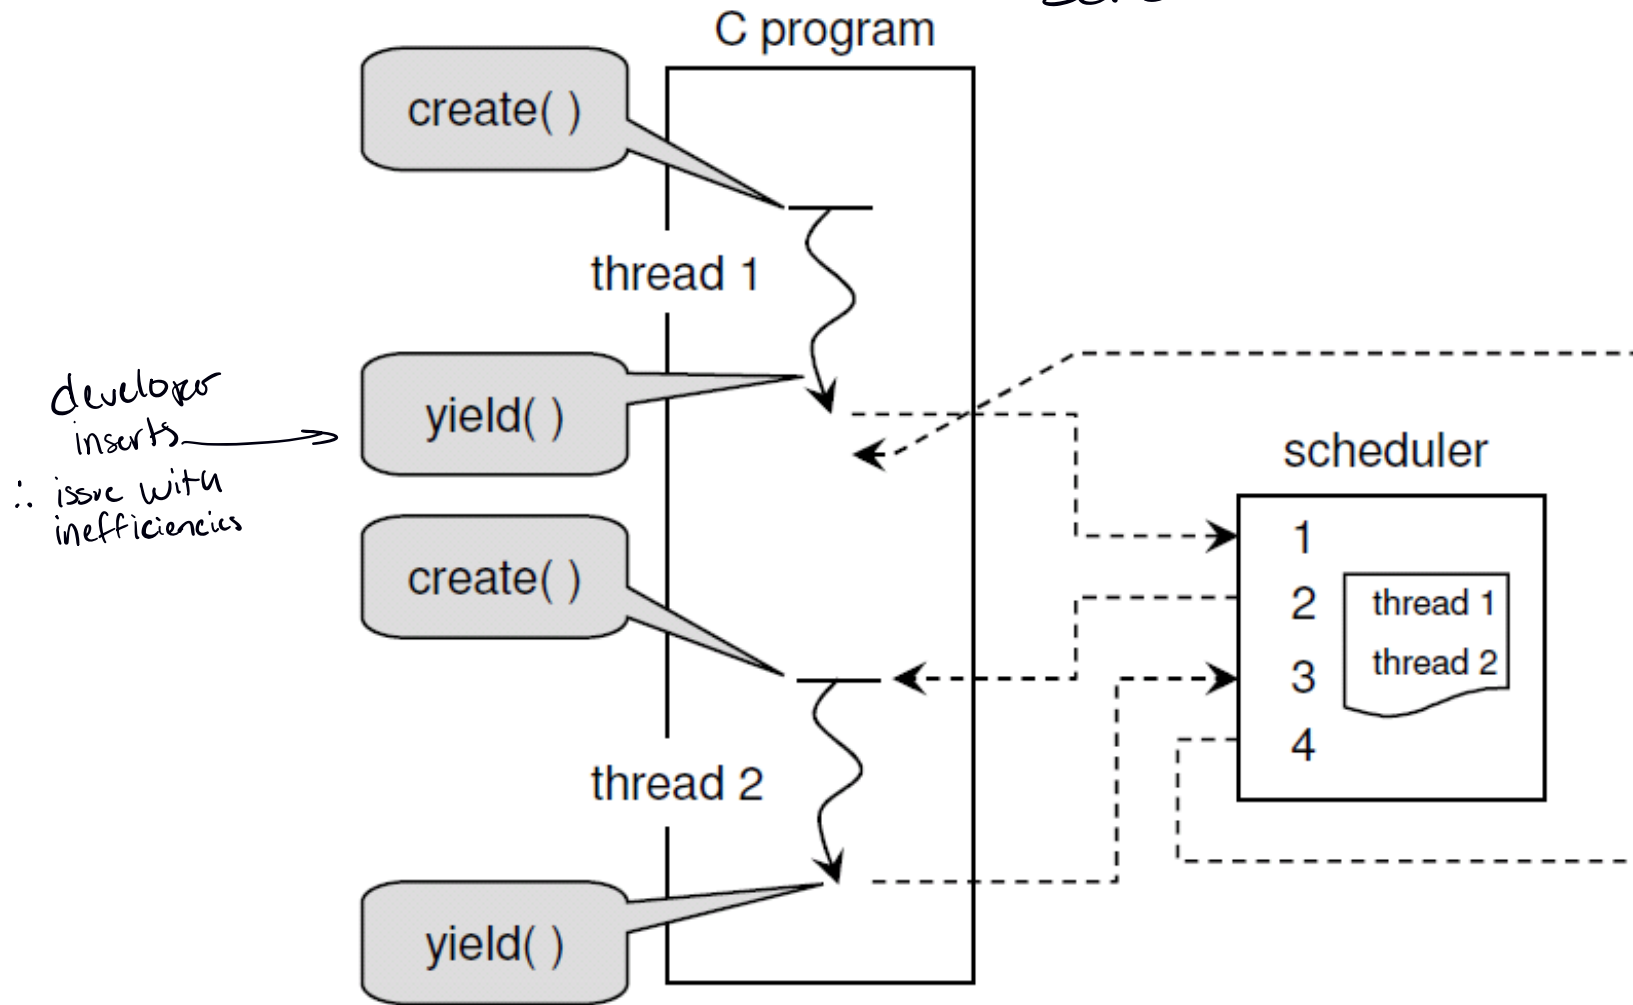
*CUDA ∴ some easier to parallelize
GPU might be slower than CPU (all unique)
but if components can be blocked and
identical, GPU faster*

Conclusions about concurrency

- If you develop a concurrent specification, you will be able to make optimal use of the underlying parallel hardware.
- In contrast, if you restrict yourself to a sequential specification from the start, it will be much harder to make use of underlying parallel hardware.
- Do not settle for a sequential language (such as C) as a universal specification mechanism. The use of C is excellent to make an executable, functional model of what you want to build.
 - But there exist concurrent specification mechanisms (we *will discuss later*) that may be better suited for parallel implementation than C.

Cooperative Multithreading

multithreading
two separate tasks, yield allows
scheduler to switch to other operation



Why do we need a scheduler?

based on yield or other things :: other policies

- Because we want to make sure that each thread gets a chance to run. The threads themselves do not have to know what other threads are active in the system.
- Because we want to have *yield points* at any point in a C program, even in the middle of a function. C does not allow to ‘suspend’ a function and then later resume it.
- Scheduling policies:
 - Round-Robin
 - Priority
 - More details coming up later in the course.

Cooperative Multithreading Example

- An example cooperative multithreading library: *quickthreads*.
- quickthreads API consists of four function calls:
 - `stp_init()` initializes the theading system
 - `stp_create(stp_userf_t *F, void *G)` creates a thread that will start execution with user function F. The function will be called with a single argument G. The thread will terminate when that function completes, or when the thread aborts.
 - `stp_yield()` releases control over the thread to the scheduler.
 - `stp_abort()` terminates a thread, so that it will be no more scheduled.

Sample code using *quickthreads*

```
#include "../qt/stp.h"
#include <stdio.h>
void hello(void *null) {
    int n = 3;
    while (n-- > 0) {
        printf("hello\n");
        stp_yield();
    }
}
void world(void *null) {
    int n = 5;
    while (n-- > 0) {
        printf("world\n");
        stp_yield();
    }
}
```

```
int main(int argc, char
**argv) {
    stp_init();
    stp_create(hello, 0);
    stp_create(world, 0);
    stp_start();
    return 0;
}
```

Summary

- With hardware, we will refer to **single-clock synchronous digital circuits**, and with software we will refer to **single-thread sequential programs**.
- Hardware and software have complementary properties for many different aspects, and are in fact each others' dual in many respects.
- The opposite of concurrency is **sequential**. Concurrent specifications are preferable over sequential ones, because it is easier to convert a concurrent specification to a parallel one than it is to parallelize a sequential specification.
- There are several modeling abstraction levels for software and hardware: **transaction-level, instruction-accurate, cycle-accurate, event-driven and continuous-time**.

Reading

For those interested (optional):

- Hillis and Steele, “Data Parallel Algorithms”, 1986:
<http://cva.stanford.edu/classes/cs99s/papers/hillis-steele-data-parallel-algorithms.pdf>
- David Patterson, “The Trouble with Multicore”, 2010:
<http://spectrum.ieee.org/computing/software/the-trouble-with-multicore>

Everyone:

For next lecture:

- If you have not taken a computer organization/architecture course before:
 - Study the basic **5-stage RISC processor pipeline**
 - http://en.wikipedia.org/wiki/Instruction_pipeline
 - http://en.wikipedia.org/wiki/Classic_RISC_pipeline
 - Hennessy & Patterson Computer Organization / Computer Architecture books are great resources

In Class Exercise (getting started for HW1)

```
#include <stdlib.h>
#include <pthread.h>
```

A simple *pthread*s example

```
void *myFirstThread(void *p)
{
    sleep(1);
    printf("I am the first thread \n");
    return NULL;
}
void *mySecondThread(void *p)
{
    printf("I am the second thread \n");
    pthread_yield();
}
void *myNewThread(void *p)
{
    printf("Hello!\n");
}
```

In Class Exercise (getting started for HW1)

```
int main()
{
    pthread_t tid, tid2, tid3;

    printf("Program started...\n");

    pthread_create(&tid, NULL, myFirstThread, NULL);
    pthread_create(&tid2, NULL, mySecondThread, NULL);
    pthread_create(&tid3, NULL, myNewThread, NULL);

    pthread_join(tid, NULL);
    pthread_join(tid2, NULL);
    pthread_join(tid3, NULL);

    printf("All threads done.\n");
    exit(0);
}
```

- What is the output?
- What happens if you do not use “join”?

Compile with: `gcc hw1intro.c -lpthread -o hw1intro`


```
#include <stdlib.h>
#include <pthread.h>

void *myFirstThread(void *p)
{
    sleep(1);
    printf("I am the first thread \n");
    return NULL;
}

void *mySecondThread(void *p)
{
    printf("I am the second thread \n");
    pthread_yield();
}

void *myNewThread(void *p)
{
    printf("Hello!\n");
    //Ask user to enter an integer
    //print that integer in decimal, binary,
    //and in hexadecimal form
}
```

***in-class exercise to
prepare for HW1***



Email your modified in-class exercise code (hw1intro.c) to:

ece535submit@gmail.com

with subject line:

yourBUusername_hw1inclass